

Module 3

Modeling / Prediction

$$Y = f(X) + \varepsilon$$

Y: Dependent Variable

f(): Some function

X: Input Variables

ε : Noise

- Our task is relatively straightforward: in general we have X, we may have some or all values of Y, ε we can do nothing about; we want to learn an approximation of f().
- Why do we want to learn f()? **Prediction** and **Inference**

A simple example: College Graduation Rates

$$Y = f(X) + \varepsilon$$

Y: Dependent Variable

f(): Some function

X: Input Variables

ε : Noise

- X: Student / Faculty Ratio, or any other property of a college.
- Y: The Graduation Rate (% of frosh that go on to graduate, I assume).
- f(): Whatever we want it to be. We start with a simple line.

Linear Model

- $f(X) = b_0 + b_1 * X_1$
- b_0 -> The intercept
- b_1 -> The slope of the dependent variable X_1 .
- $Y = f(X) + \varepsilon \rightarrow Y = b_0 + b_1 * X_1 + \varepsilon$
- We learn $f(X)$ in this case through linear regression with the `lm()` function.

Data set

- Using the college dataset again:

```
college <- read.csv("College.csv")
```

```
rownames(college) <- college$X
```

```
college$X <- NULL
```

```
college <- college[college$Grad.Rate < 100,]
```


Investigate by eye

- Do we see a relationship between student / faculty ratio and grad. rate?

```
plot(college$S.F.Ratio,college$Grad.Rate)
```


Yes.

Lower S.F rates seem to correlate with higher Graduation Rates.



Learn linear model

```
lm.fit <- lm(Grad.Rate ~ S.F.Ratio,data=college)
abline(lm.fit)
print(summary(lm.fit))
```



Learn linear model

```
> summary(lm.fit)
```

Call:

```
lm(formula = Grad.Rate ~ S.F.Ratio, data = college)
```

Residuals:

Min	1Q	Median	3Q	Max
-54.006	-10.690	0.724	11.640	39.502

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	84.0830	2.1268	39.536	<2e-16 ***
S.F.Ratio	-1.3583	0.1454	-9.345	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.86 on 764 degrees of freedom

Multiple R-squared: 0.1026, Adjusted R-squared: 0.1014

F-statistic: 87.32 on 1 and 764 DF, p-value: < 2.2e-16

Prune outlier

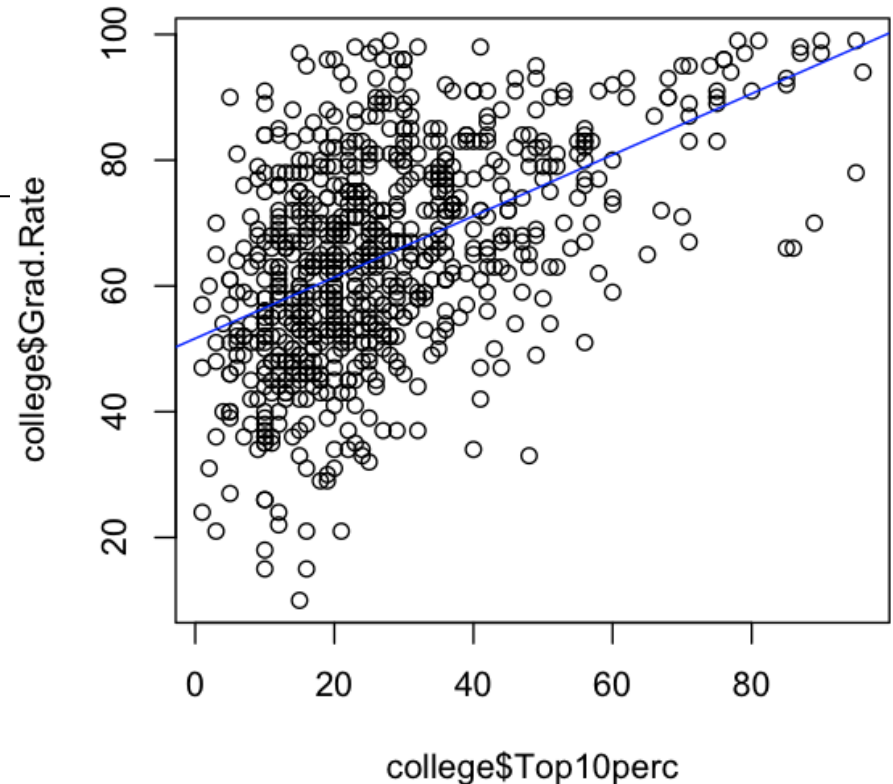
```
college.s <- college[college$S.F.Ratio < 30,]  
plot(college.s$S.F.Ratio,college.s$Grad.Rate)  
lm.fit.s <- lm(Grad.Rate ~ S.F.Ratio,data=college.s)  
abline(lm.fit.s)  
print(summary(lm.fit.s))  ## basically the same
```

- Should see that the outlier was having very little effect.

What about % students from the top 10% of their graduating HS class?

```
plot(college$Top10perc,college$Grad.Rate)  
lm.fit <- lm(Grad.Rate ~ Top10perc,data=college)  
abline(lm.fit,col='blue')  
print(summary(lm.fit))
```

A pretty good relationship:



What model have we learned?

```
coef(lm.fit)
```

- (Intercept) Top10perc
- 51.6430252 0.4866471

```
confint(lm.fit)
```

- So, our model ($Y=f(X) + \epsilon$) now is:

$$\text{Grad.Rate} = 51.6 + .48 * \text{Top10perc} + \epsilon$$

What if we had 4 new colleges we hadn't seen before. Could we guess their graduation rate?

```
newcolleges <- data.frame(  
  CollegeName=c("MattU", "PavoTech", "ApoorvaCollege", "SheamusInstitute"),  
  Top10perc=c(50, 60, 99, 5)  
)
```

```
rownames(newcolleges) <- newcolleges$CollegeName
```

```
predict(lm.fit, newdata=newcolleges)
```

•	MattU	PavoTech	ApoorvaCollege	SheamusInstitute
	75.97538	80.84185	99.82108	54.07626

```
predict(lm.fit, newdata=newcolleges, interval="prediction")
```

•		fit	lwr	upr
•	MattU	75.97538	47.49454	104.45621
•	PavoTech	80.84185	52.32649	109.35720
•	ApoorvaCollege	99.82108	71.05317	128.58900
•	SheamusInstitute	54.07626	25.59637	82.55615

Diagnostic plot




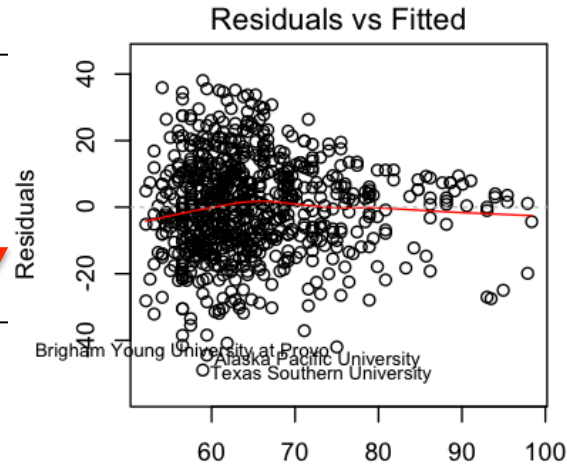
```
par(mfrow=c(2,2))  
plot(lm.fit)
```

Diagnostic plot

```
par(mfrow=c(2,2))  
plot(lm.fit)
```

How wrong we were

Our guess



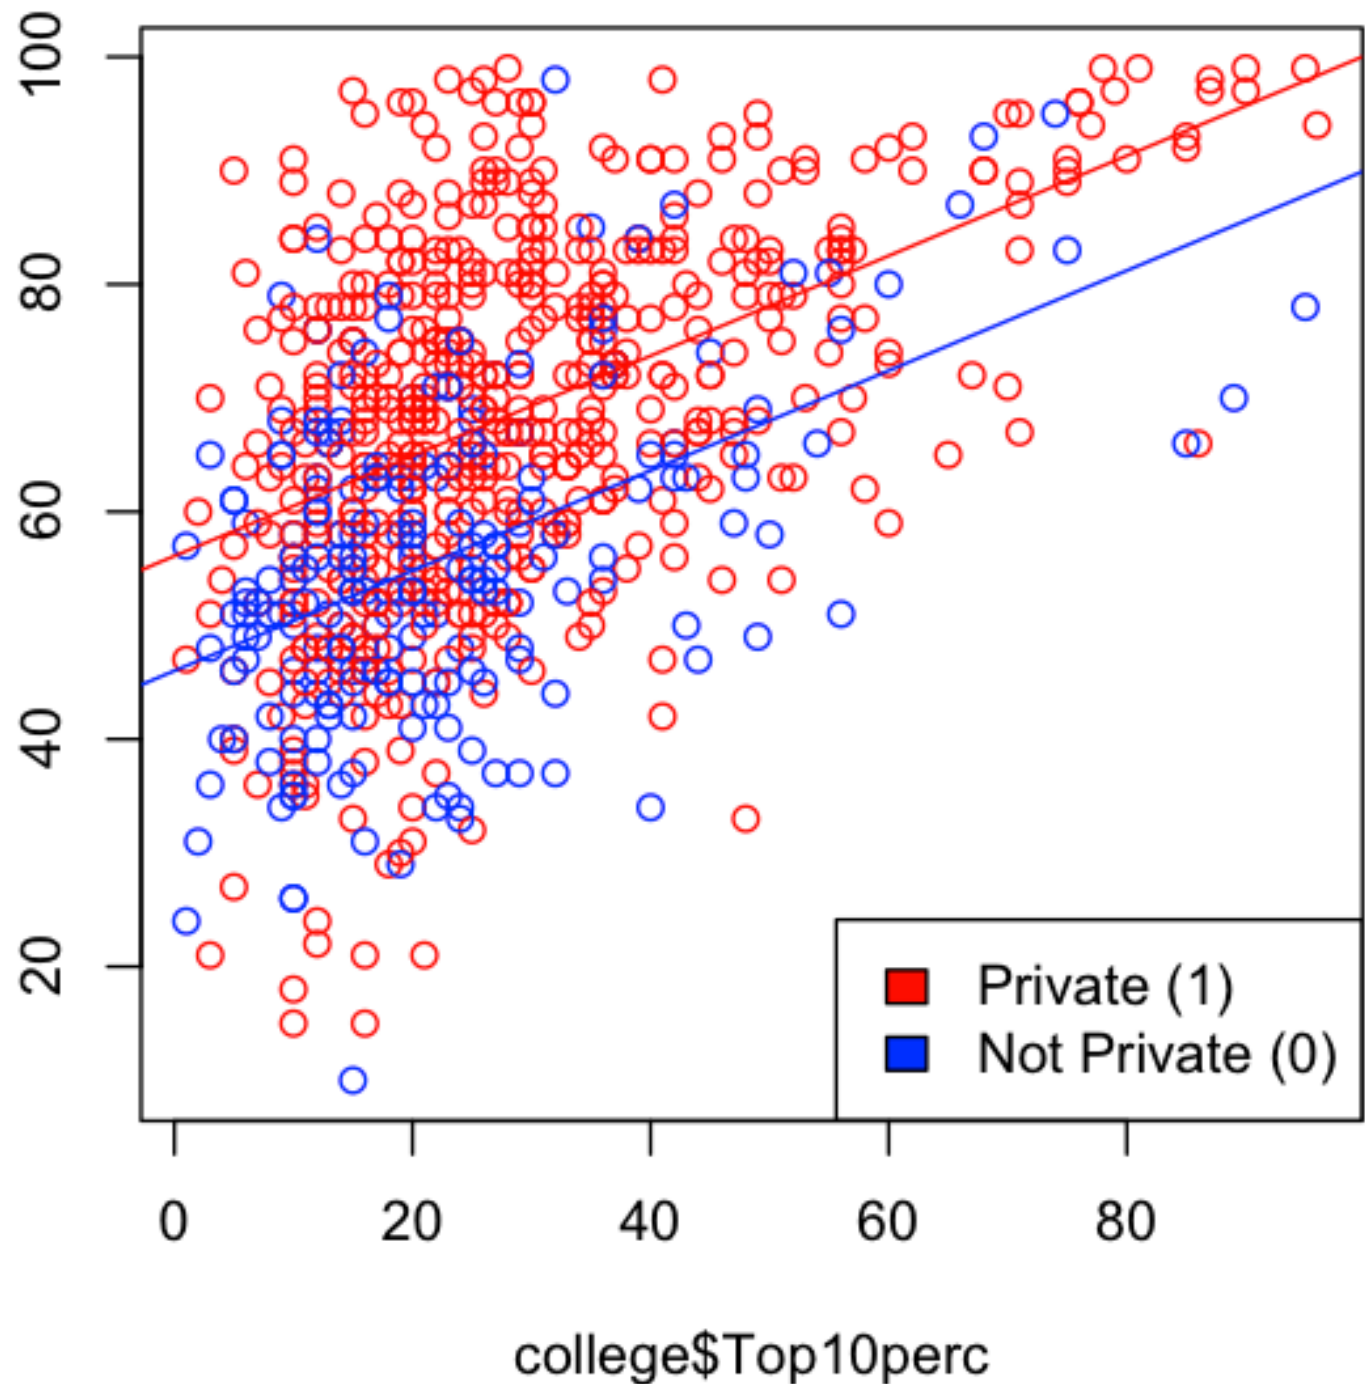
Multiple regression

- $f(X) \rightarrow b_0 + b_1 * X_1 + b_2 * X_2$
- Learning over multiple features simultaneously.
- Keep X_1 as Top10perc, make X_2 Private
- Private an indicator variable, so value of X_2 will be $\{0,1\}$, so really just learning a separate intercept:
 - $\text{Grad.Rate} = b_0 + b_1 * (\text{Top10Perc}) + b_2 * \{0,1\}$
 - If not private:
 - $\text{Grad.Rate} = b_0 + b_1 * (\text{Top10perc})$
 - If private:
 - $\text{Grad.Rate} = b_0 + b_1 * (\text{Top10perc}) + b_2$

Multiple regression

```
lm.fit.mult <- lm(Grad.Rate ~ Top10perc + Private,data=college)
print(summary(lm.fit.mult))
coef(lm.fit.mult)
par(mfrow=c(1,1))
plot(college$Top10perc,college$Grad.Rate,col=ifelse(college
$Private=="Yes",'red','blue'))
abline(lm.fit.mult,col='blue')
abline(a=coef(lm.fit.mult)[1] + coef(lm.fit.mult)[3], b=
coef(lm.fit.mult)[2],col='red')
legend('bottomright',legend=c("Private (1)","Not Private
(0)"),fill=c("red","blue"))
```

```
lm.fit.mult <-
print(summary
coef(lm.fit.n
par(mfrow=c(1
plot(college$
$Private=="Ye
abline(lm.fit
abline(a=coef
coef(lm.fit.n
legend('bottc
(0)'),fill=c(
```




Multiple regression

```
## what about just using every variable?  
## The code for that is "~."  
lm.fit.all <- lm(Grad.Rate ~ ., data=college)  
print(summary(lm.fit.all))
```

- As you can see from the summary, the fit appears to be good, as there's a high R^2 : 0.48 vs 0.31 for our last model.
- However, this model is almost certainly overfit, given the number of free parameters vs the number of colleges.
- Additionally, the significances and weights are difficult or impossible to interpret since most of the input variables are highly correlated.

Logistic Regression

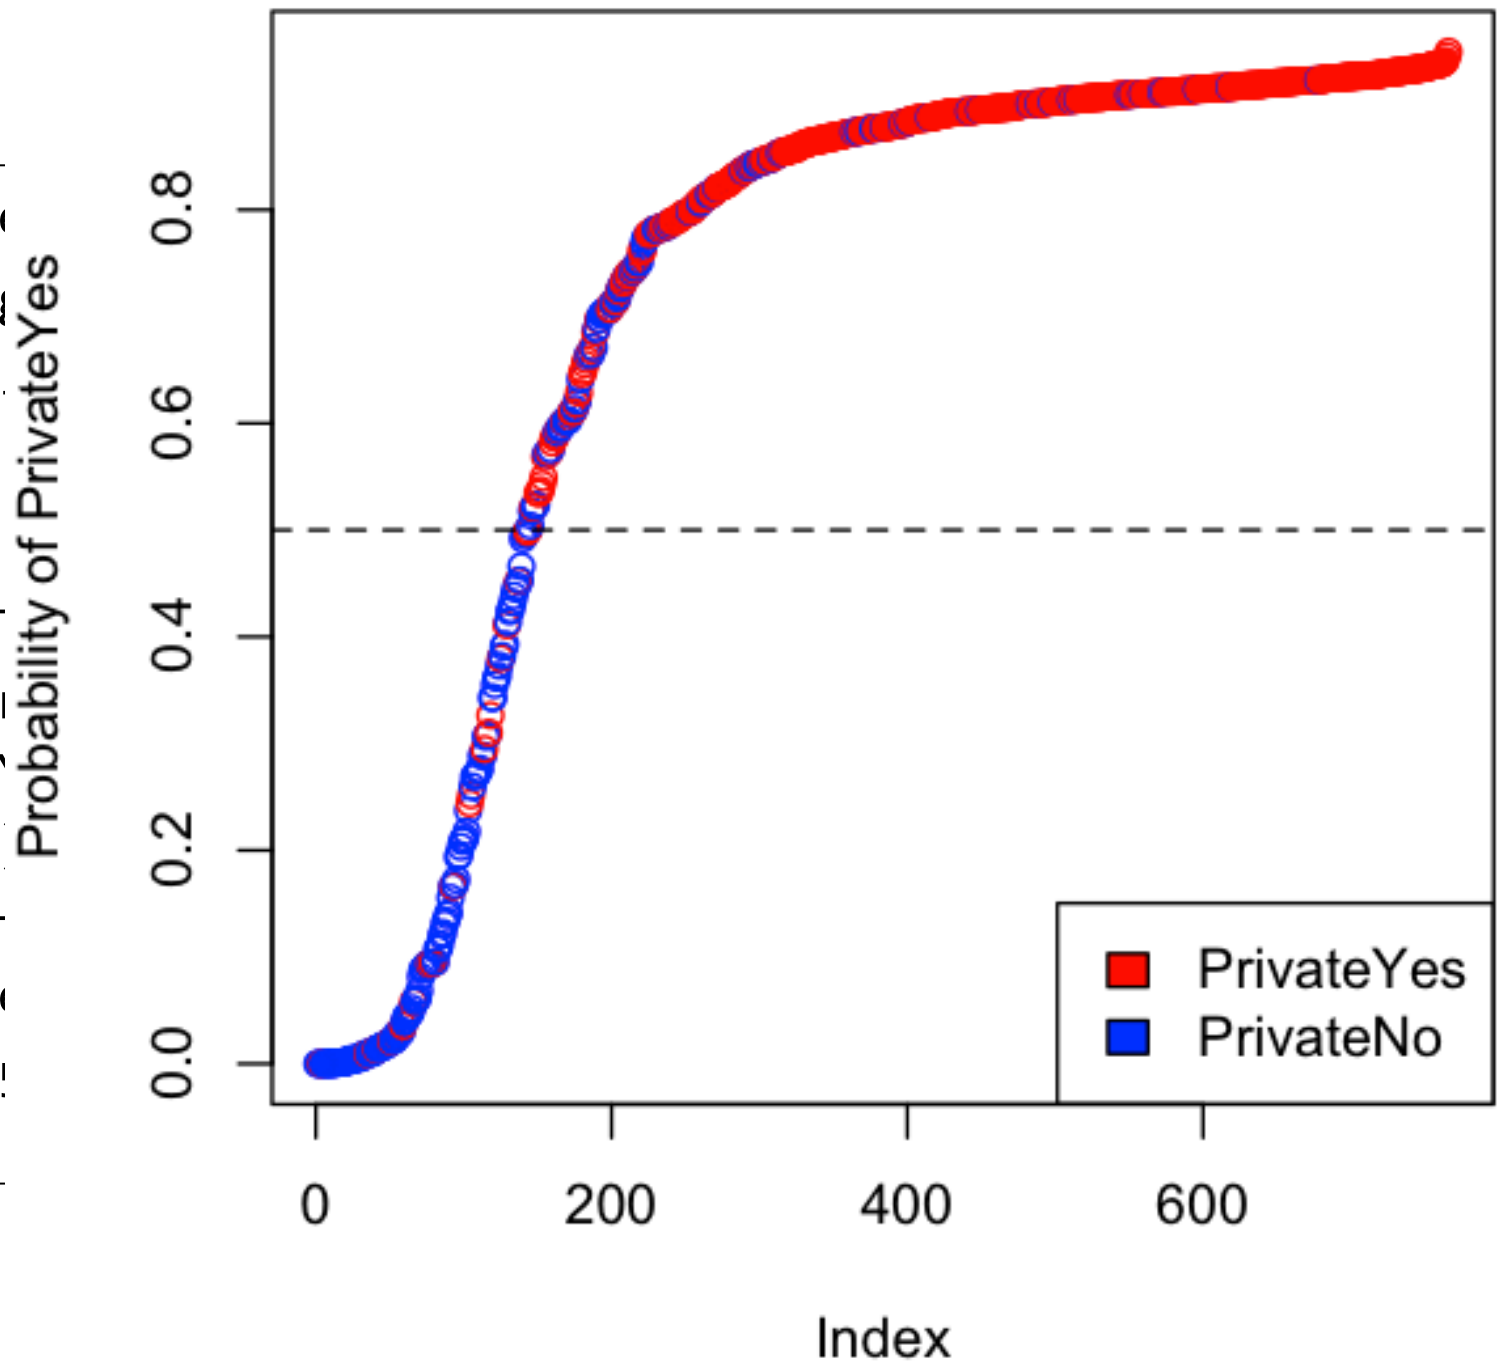
- Goal: Make a true / false classifier.
- Replace $f(X)$ with a logistic function.



Logistic regression

```
college$AcceptanceRate <- college$Accept / college$Apps  
glm.fit <- glm(Private ~ AcceptanceRate +  
F.Undergrad,data=college,family='binomial')  
predict(glm.fit,type='response')  
## plot the results  
p <- predict(glm.fit,type='response')  
plot(p[order(p)],col=ifelse(college  
$Private=="Yes",'red','blue')[order(p)],ylab="Probability of  
PrivateYes")  
legend('bottomright',legend=c("PrivateYes","PrivateNo"),fill=c  
("red","blue"))  
abline(h=0.5,lty=2)
```

```
college$Acco
glm.fit <- g
F.Undergrad
predict(glm
## plot the
p <- predic
plot(p[order
$Private=="
PrivateYes"
legend('bot
("red", "blue
abline(h=0.5
```



Logistic regression

```
table(PredictedPrivate=p > .5, isPrivate=college  
$Private)
```

	isPrivate	
PredictedPrivate	No	Yes
FALSE	124	19
TRUE	87	536

- Pretty good accuracy!
- But, we're over-determined, since we're testing on the same data that we learned our regression on. This means the above accuracy is not to be trusted.
- A better way: cross validation.

Cross Validation

- An unbiased way to see if the classifier you're learning would be accurate on new data, for example a new set of colleges you'd never seen before.
- 10-fold cross validation:
 - Learn the classifier on 90% of the data, test on remaining 10%.
 - Repeat 10 times, until you've tested on every record.

Logistic regression

```
college$CV.index <- rep(1:10,length=nrow(college))
for(i in 1:10) {
  in.fold <- college$CV.index != i
  glm.fit.cv <- glm(Private ~ AcceptanceRate + F.Undergrad,
data=college[in.fold,],family='binomial')
  p <- predict(glm.fit.cv,type='response',newdata=college[!in.fold,])
  college[!in.fold,'prediction'] <- p
}
table(PredictedPrivate=college$prediction > .5, isPrivate=college$Private)
```

	isPrivate	
PredictedPrivate	No	Yes
FALSE	122	20
TRUE	89	535


- Wow! Still good accuracy!

ROC Curve

```
install.packages("ROCR")  
library("ROCR")  
pred <- prediction(college$prediction,college$Private=="Yes")  
perf <- performance(pred,measure='tpr',x.measure='fpr')  
perf.auc <- performance(pred,measure='auc')  
plot(perf)  
abline(a=0,b=1,lty=3)  
text(.5,.2,paste("AUC:",formatC(perf.auc@y.values[[1]])))
```

- A standard way to judge classifier sensitivity and specificity is to use a Receiver Operator Characteristic (ROC) curve and measure the area under the curve (AUC).

ROC Curve



- Great ROC curve!
- Along the diagonal signals random data.
- Up the y axis signals perfect classification.
- .898 is a very high AUC!